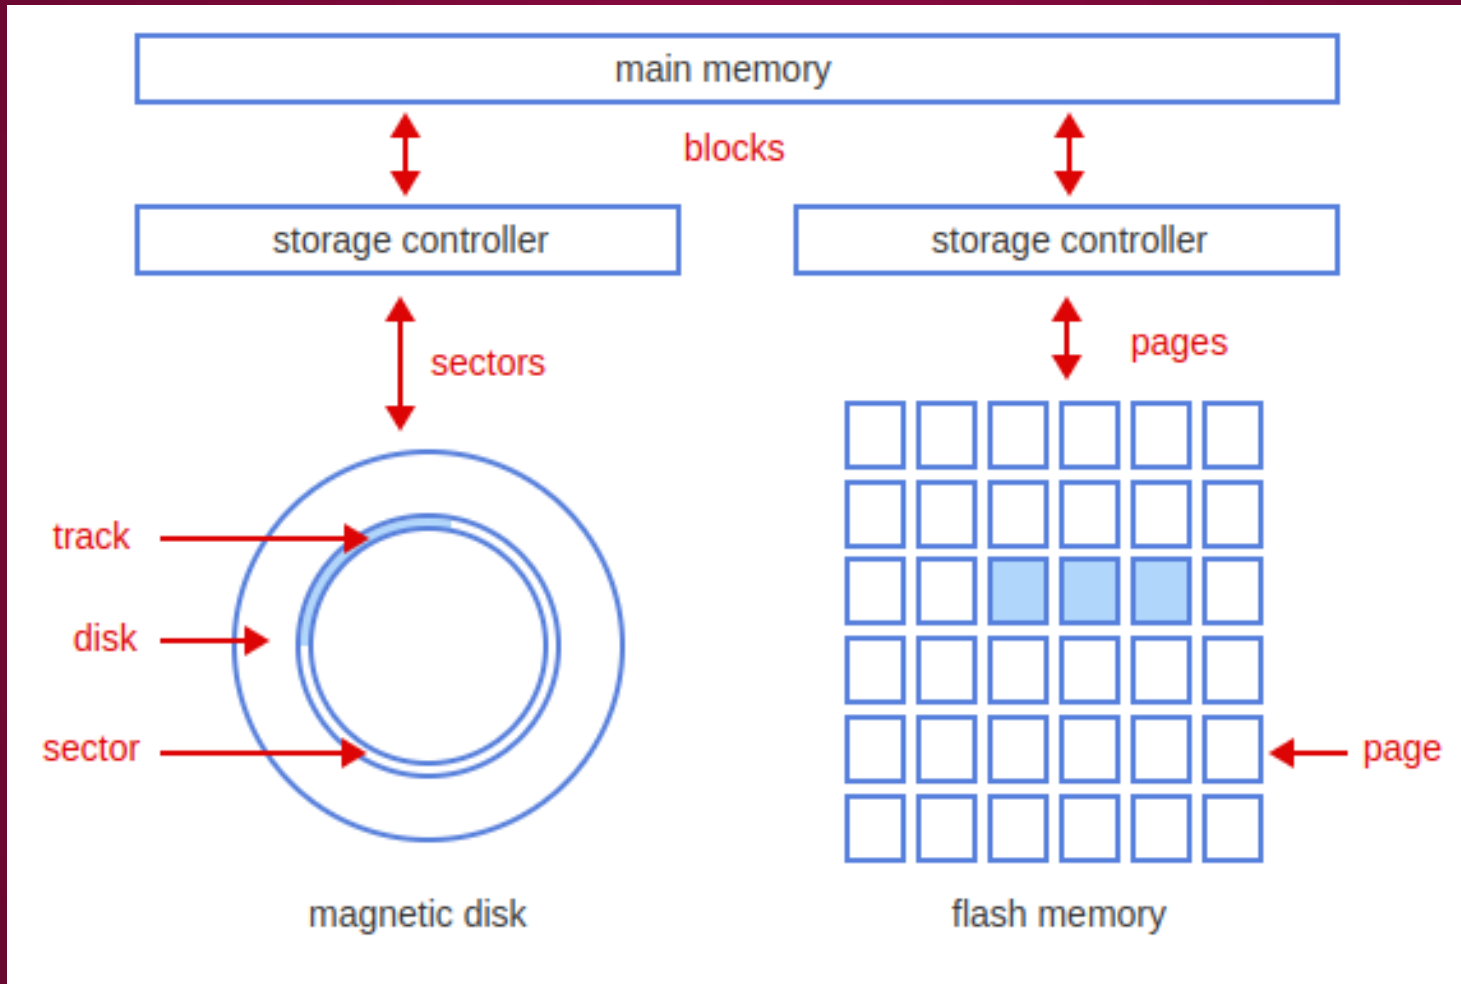


# Physical Design

# Working in the background

- Decisions affect performance but never affect query results

# Storage Media



# Data Organization

- Row oriented
  - Better for transactional applications
- Column oriented
  - Better for analytical applications

# Table Structure

- Heap - No order is imposed
  - Fast write, slow read
- Sorted table - Rows are sorted by a column
  - Fast read, slow write
- Hash table - Data is grouped together based on hash
  - Fast for individual reads/writes, slow for bulk
- Table cluster - Multiple tables interleaved
  - Specialized for a particular join, slow for most other operations

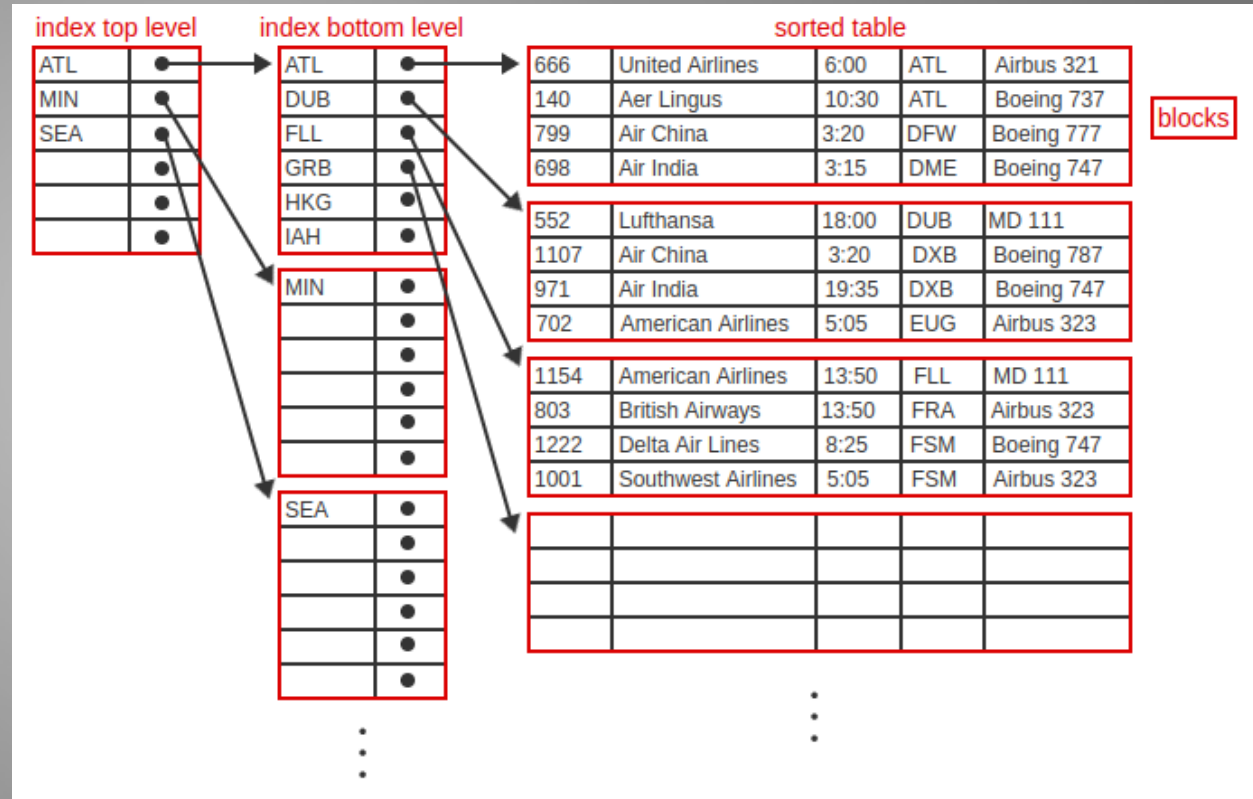
# Indexes

- Addresses the main weakness of a heap table
- One column (along with a pointer) is sorted
- Allow for the possibility of a binary search
- Multiple indexes?
- Dense vs Sparse

Socratica Video

# Multi-level Indexes

- B+tree index
  - Values are repeated in lower level indexes
- B-tree index
  - Values are not repeated in lower level indexes



# Binary Search - Find 69951757

4055505  
6995630  
9086405  
9270322  
9362777  
10484970  
11848518  
12622330  
13440190  
13689200  
14082277  
15347211  
15419461  
16483496  
18521572  
19870627  
20138740  
20406038  
21666247  
23408187

25419539  
28001595  
28867473  
29515064  
30139471  
32853477  
32964207  
33657355  
33671279  
34096808  
35700734  
36312160  
39128058  
40095578  
42240036  
42942475  
43247718  
43541621  
49147525  
51034515

51615167  
54182046  
54458130  
55967518  
56678759  
58751875  
60528887  
61862972  
62360432  
63236935  
63420879  
69951757  
70295119  
71129543  
71872482  
72060865  
74782400  
76091035  
76327014  
76337118

77100229  
77193195  
79727269  
80468817  
80769463  
86354069  
89028599  
89965459  
92439732  
92957685  
93201595  
93510255  
93696884  
95438865  
95628831  
96039351  
96349184  
99070162  
99569182  
99945508



# Storage Engine

- InnoDB - Transactional DB (Default)
  - Heap or sorted tables
  - B+tree indexes
- MyISAM - Analytical applications with limited updates
  - B+tree indexes
- MEMORY - all data is stored in memory
  - B+tree or hash indexes

# SQL for working with indexes

- `CREATE INDEX IndexName on TableName (Column1, Column2, etc);`
- `DROP INDEX IndexName on TableName`
- `SHOW INDEX FROM TableName;`

# EXPLAIN Statement

- Shows how an CRUD statement will be executed
- Help with deciding on an index
- Identify slow queries

```
EXPLAIN
  SELECT FlightNumber
  FROM Flight
 WHERE AirportCode IN
    (SELECT AirportCode
     FROM Airport
     WHERE Country = "Cuba");
```